

Sprint Report

Unit code: COS40005

Unit name: Computing Technology Project B

Submission date: 23 Aug 2026

Contribution for: Sprint 4

Table of Contents

Acknowledgment of Country	2
Contribution summary (Sprint 4)	3
1. Sprint plan	4
Sprint Goal.....	4
Scrum Master	4
Planned Artifacts	4
2. Risk management plan	5
3. Sprint Progress.....	7
Lakshit Bansal	8
Ved Jay Makhijani.....	8
Kumarawadu Dinith Thejana.....	10
Kaveesha Dulanjana Dharmadasa (105271678)	14
Minh Hoang Duong	17
4. Sprint review (Critical review of the product).....	19
Demonstration Session	19
Client Feedback	19
Empirical Performance Summary	19
Critical Analysis of Progress and Results	20
How the System is Safe	21
Why the Reported Error Rate is 0%	22
Critical Analysis of the Process.....	22
5. Retrospect (Critical review of the process)	24
6. Lessons learned (Critical review of Sprint 4 experience and future plan)	25
Recommendations for the next sprint	25
7. Appendix A: Pseudocode of Key Pipeline Components	26
A.1 Pipeline Orchestrator (main.py)	26
A.2 LLM Instruction Parser (schema-first)	26
A.3 Task Planner (rule-based decomposition)	26
A.4 Vision / Scene Representation (with graceful fallback)	27
A.5 Execution Layer (boundary-checked)	27
Appendix B: References	27

ACKNOWLEDGMENT OF COUNTRY

Ved Jay Makhijani

I, Ved Jay Makhijani, respectfully acknowledge the Wurundjeri People of the Kulin Nation, who are the Traditional Owners of the land on which Swinburne's Australian campuses are located in Melbourne's east and outer-east, and pay our respect to their Elders past, present and emerging.

Minh Hoang Duong

I, Minh Hoang Duong, respectfully acknowledge the Wurundjeri People of the Kulin Nation, who are the Traditional Owners of the land on which Swinburne's Australian campuses are located in Melbourne's east and outer-east, and pay our respect to their Elders past, present and emerging.

Kumarawadu Dinith Thejana

I, Kumarawadu Dinith Thejana, respectfully acknowledge the Wurundjeri People of the Kulin Nation, who are the Traditional Owners of the land on which Swinburne's Australian campuses are located in Melbourne's east and outer-east, and pay our respect to their Elders past, present and emerging.

Kaveesha Dulanjana Dharmadasa

I, Kaveesha Dulanjana Dharmadasa, respectfully acknowledge the Wurundjeri People of the Kulin Nation, who are the Traditional Owners of the land on which Swinburne's Australian campuses are located in Melbourne's east and outer-east, and pay our respect to their Elders past, present and emerging.

Lakshit Bansal

I, Lakshit Bansal, respectfully acknowledge the Wurundjeri People of the Kulin Nation, who are the Traditional Owners of the land on which Swinburne's Australian campuses are located in Melbourne's east and outer-east, and pay our respect to their Elders past, present and emerging.

CONTRIBUTION SUMMARY (SPRINT 4)

* Attendance fields were left as Yes* where the source table was blank; verify these against the meeting records before final submission.

Team Member	Tasks completed	Initiative	Team meetings	Supervisor meetings	Client meetings	Communication	Status updates	Respect
Lakshit Bansal	Yes – Evaluation expansion, latency benchmarking and Scrum Master responsibilities completed.	Yes – Proactively expanded evaluation coverage and introduced end-to-end latency benchmarking.	Yes*	Yes*	Yes*	Yes – Active project communication.	Yes – Coordinated sprint progress and blockers.	Yes
Ved Jay Makhijani	Yes – Evaluation, benchmarking and final report sections completed.	Yes – Resolved compatibility issues and strengthened evaluation methodology.	Yes*	Yes*	Yes*	Yes – Active team communication.	Yes – Shared evaluation and report progress.	Yes
Minh Hoang Duong	Yes – Repository cleanup, refactoring and integrity audit completed.	Yes – Proactively identified and fixed structural, path and configuration issues.	Yes*	Yes*	Yes*	Yes – Active technical communication.	Yes – Shared refactoring and integration updates.	Yes
Kaveesha Dulanjana Dharmadasa	Yes – Discussion and conclusion sections completed.	Yes – Performed source-to-report consistency and evidence checks.	Yes	Yes	Yes	Yes – Active documentation communication.	Yes – Shared report updates.	Yes
Dinith Thejana	Yes – YOLOv8 integration, PyBullet execution and end-to-end testing completed.	Yes – Generated dataset, fine-tuned YOLOv8 and resolved execution issues.	Yes	Yes	Yes	Yes – Active technical communication.	Yes – Shared vision and integration results.	Yes

1. SPRINT PLAN

Sprint Goal

The main goal of Sprint 4 is to progress the project from the Semester 1 Alpha prototype towards a Beta release. The existing end-to-end pipeline has already been established, connecting natural language instruction, visual scene information, task planning and robot execution. Therefore, this sprint focuses on improving the reliability, stability and validation of the integrated system rather than rebuilding the core pipeline.

The main objectives are to stabilise the live vision and scene representation components, validate the interface between vision and task planning, improve execution reliability and error handling, and perform end-to-end testing using different task and scene conditions.

Scrum Master

Scrum Master: Lakshit Bansal

The Scrum Master will coordinate sprint activities, monitor progress against the sprint backlog, identify blockers and ensure that communication between team members remains consistent throughout the sprint.

Planned Artifacts

The planned Sprint 4 artifacts include:

- Updated live scene representation and vision interface.
- Updated execution and error-handling components.
- End-to-end integration and regression tests.
- Beta testing results and evidence.
- Updated risk register.
- Updated technical documentation and project evidence.

The focus of these artifacts is to improve the reliability and reproducibility of the existing integrated system and prepare it for the next stage of development.

The project plan originally identified the Semester 2 Beta release as requiring full system integration, improved accuracy, testing and bug fixing. The current sprint therefore builds on the completed Alpha functionality rather than repeating the integration work already achieved.

2. RISK MANAGEMENT PLAN

The Sprint 4 risk management plan continues the risks identified during Semester 1 and adds risks associated with the transition from the Alpha prototype to the Beta release. During Semester 1, the main technical risks included PyBullet compatibility, external LLM dependencies, integration issues and the reliability of AI-generated commands

Risk	Likelihood	Impact	Mitigation
PyBullet/environment compatibility	Medium	High	Maintain a documented environment setup, test on supported systems and keep simulation configuration consistent across team members.
Vision detection errors	Medium	High	Validate live detections against the known simulation scene and use the existing fallback mechanisms when required.
Vision-to-planner schema mismatch	Medium	High	Maintain a clearly defined scene contract and test the live scene output before it reaches the planner.
Incorrect coordinates causing unsafe movement	Medium	High	Apply workspace boundary validation and verify coordinates between vision, planning and execution stages.
LLM incorrect or ambiguous output	Medium	High	Use structured schemas, validation, confidence handling and reject invalid instructions before execution.
External API availability or latency	Medium	Medium	Maintain multiple supported LLM backends and retain the local HuggingFace backend as an alternative.
Integration regression after code changes	Medium	High	Run automated regression and end-to-end tests after significant changes and review changes before merging.

Insufficient testing of failure cases	Medium	High	Include invalid objects, ambiguous instructions, execution failures and different scene configurations in Beta testing.
Dependency and package compatibility issues	Medium	Medium	Keep requirements documented and test the project in a clean virtual environment before the Beta demonstration.
Privacy of industrial instructions	Low/Medium	High	Avoid using real confidential industrial data and consider local LLM inference for future deployment involving sensitive instructions.

The project continues to recognise that external LLM providers may introduce privacy concerns for real industrial instructions. This was already identified in Semester 1, where the team proposed local model inference as a mitigation.

3. SPRINT PROGRESS

Sprint 4 has focused on stabilising the existing Alpha system and preparing it for Beta-level testing. The project entered the sprint with the major pipeline components already implemented, including the LLM parser, visual perception, task planner, execution layer and PyBullet simulation.

The current codebase contains a production `main.py` pipeline which connects the major stages of the system. The pipeline accepts a natural language instruction, obtains the current scene from the live simulation vision stack, generates a task plan and passes the resulting commands to the executor. This confirms that Sprint 4 is focused on improving and validating the existing integration rather than creating the end-to-end pipeline from scratch.

The live scene adapter in `vision_backend/scene_representation.py` has also been developed to provide a consistent planner-facing scene structure containing object labels and 3D positions. The adapter obtains scene information from the simulation's live vision system and validates that expected workspace objects are present.

The team has also continued testing the integration between the vision, planning and execution layers. The repository contains integration tests covering spatial relationships, task planning and execution behaviour, together with tests for the real vision adapter.

During the sprint, development has also focused on identifying runtime and environment issues that can affect execution of the complete pipeline. These issues are being addressed as part of the Beta stabilisation process.

Relevant evidence for this section includes:

- Jira Sprint 4 task board.
- GitHub commits from each team member.
- Live scene representation output.
- End-to-end pipeline terminal output.
- PyBullet robot execution screenshots.
- Automated test results.
- Evidence of successful and failed task validation.

The work builds directly on the previous sprint, where the team established the planning, execution and simulation architecture and identified real-world integration and testing as important areas for further development

Individual Contributions

Lakshit Bansal

During Sprint 4, my primary contribution focused on extending the evaluation of the embodied task planning system and measuring its end-to-end performance. I was responsible for two evaluation tasks: expanding the existing evaluation dataset and benchmarking the latency of the complete pipeline.

Evaluation Dataset Expansion (EVAL-8 / PITPP-81):

I worked on expanding the existing evaluation dataset from 25 test cases to more than 35 cases. The additional test cases focus on ARENA2036-specific industrial instructions to make the evaluation more representative of the project's intended industrial task-planning context. The expanded dataset is intended to provide a broader basis for comparing model performance and identifying weaknesses that may not have been captured by the original 25 test cases.

End-to-End Latency Benchmarking (EVAL-9 / PITPP-82):

I also worked on benchmarking the end-to-end latency of the complete system across 10 representative instructions. The benchmark measures the time taken by the system from receiving an instruction through processing and task execution. Mean latency and 95th-percentile (P95) latency are being calculated to provide a more representative measure of typical and worst-case system response time.

These evaluation tasks extend the model evaluation completed in the previous semester by increasing the diversity of the test dataset and introducing end-to-end performance measurements. The results will be used to assess the system's reliability and suitability for real-time industrial task planning.

Ved Jay Makhijani

During Sprint 4 I completed five independent tickets spanning evaluation, benchmarking, and final report writing. Each deliverable below presents the action taken and its direct output together.

EVAL-7 — Full GPT-4o Pipeline Evaluation

Fixed all Python 3.14 f-string incompatibilities in eval_report.py — nested format specifiers such as `<{width}`, `>{width}`, and `^{width}` were replaced with equivalent string method calls (`.ljust()`, `.rjust()`, `.center()`). Also fixed dictionary key access inside f-strings by extracting values to local variables before the print call. Successfully ran the full evaluation pipeline across all 25 test cases using GPT-4o and the rule-based baseline via `python -m llm_backend.LLM_eval.eval_report --models openai --export`.

Output: Full metrics comparison table showing GPT-4o vs baseline across parse success rate, instruction accuracy, action accuracy, object accuracy, destination accuracy, spatial accuracy, avg latency, and error rate. Results exported to evaluation_metrics.csv and evaluation_results.json. Terminal screenshot of the full metrics table submitted as evidence.

EVAL-8 — Extended Evaluation Dataset to 35 Cases

Extended the evaluation dataset from 25 to 35 labelled test cases in llm_backend/LLM_eval/test_cases.py. Added 10 new cases covering ARENA2036-style industrial language (transfer the component to assembly station A), harder spatial configurations (place the red block directly behind the workstation), compound synonym instructions (grab and deliver the green block to the right tray), and adversarial edge cases (move the object — ambiguous, no object specified). Each new case includes ground truth labels for expected action, object target, destination, spatial relation, and confidence level. Baseline rerun confirms correct coverage across all categories.

Output: test_cases.py committed with 35 entries. Baseline evaluation terminal showing 35 test cases and updated category breakdown. GitHub commit submitted as evidence.

EVAL-9 — End-to-End Latency Benchmark

Built benchmark.py to measure wall-clock GPT-4o parsing latency across 10 instructions spanning all instruction categories. The script calls parse_instruction() for each instruction, records elapsed time, sorts results, and computes mean, p95, min, and max. Ran against the GPT-4o backend from Melbourne, Australia on a consumer internet connection.

Results: Mean 1,433ms | P95 2,461ms | Min 1,035ms | Max 2,461ms. All 10 instructions parsed successfully on the first attempt with no retries. No strong correlation between instruction complexity and latency — the simplest instruction (locate the green block) returned in 1,136ms while a multi-step spatial instruction returned in 1,382ms, indicating that variance is dominated by network conditions rather than model processing time.

Output: Terminal screenshot of benchmark.py output showing per-instruction latency and summary statistics. Numbers used directly in final report Section 4.4.

REP-2 — Methodology Section of Final Report

Wrote the complete methodology section of the final report covering Sections 3.1 through 3.6. Section 3.1 describes the 5-stage pipeline architecture and three core design principles (schema-first, stub-based integration, dependency-free testing). Section 3.2 covers the LLM instruction parser including schema enforcement via Pydantic v2, prompt design with six few-shot examples, edge case handling, and model portability across GPT-4o, Gemini, and DeepSeek. Section 3.3 covers the task planner including the rule-based design decision justification, action decomposition sequences, spatial relation handling via SPATIAL_OFFSETS, and multi-step

support via `plan_multi_step()`. Section 3.4 describes the execution layer including MockRobot and RealRobot architecture — explicitly noting MockRobot is not a trivial stub and maintains real internal state. Section 3.5 covers vision integration including the graceful fallback mechanism. Section 3.6 covers the evaluation methodology including the 35-case dataset design, baseline parser rationale citing HELM (Liang et al. 2022), and all 10 metrics defined and justified. Approximately 2,400 words.

Output: Final report document showing Sections 3.1–3.6 drafted. Screenshot of report document open submitted as evidence.

REP-3 — Evaluation and Results Section of Final Report

Wrote the complete evaluation and results section covering Sections 4.1 through 4.5 using real numbers from EVAL-7 and EVAL-9. Section 4.1 presents the overall model comparison metrics table — GPT-4o vs baseline across all 10 metrics. Section 4.2 provides category-level analysis for each of the six categories, with particular focus on the four most significant findings: ambiguous (0% baseline vs real GPT-4o result), multi-step (33.3% vs real), spatial (60% vs real), and simple (60% vs 100%). Section 4.3 makes the research contribution argument — that the accuracy gap on spatial, ambiguous, and multi-step categories quantifies meaningful LLM benefit over keyword matching, citing SayCan (Ahn et al. 2022) and PaLM-E (Driess et al. 2023). Section 4.4 presents the latency analysis using real benchmark numbers. Section 4.5 describes the pipeline execution results and the tracker audit trail as a safety-relevant feature. Approximately 1,350 words plus tables.

Output: Final report document showing Sections 4.1–4.5 drafted with real evaluation numbers. Screenshot submitted as evidence.

Kumarawadu Dinith Thejana

During Sprint 4, my main contribution focused on completing and testing the real vision, PyBullet simulation and end-to-end integration of the embodied task planning system. My work focused on ensuring that the complete pipeline could receive a natural-language instruction, detect objects from the simulated workspace, generate a task plan and execute the task using the KUKA robot.

VIS-FINAL — Complete Real Vision Integration:

I have been working on the completion of the YOLOv8 vision integration in the simulated workspace. The existing COCO pre-trained YOLOv8 did not detect the blocks, trays, and workstations that were custom to this project. Therefore, a synthetic dataset was generated in the PyBullet environment based on the images rendered and their segmentation information. In total, the dataset included 300 images with 2,100 object labels in seven workspace classes. The YOLOv8n model was fine-tuned with this dataset, and the detection of all seven workspace

objects was successfully verified with direct vision. The fine-tuned model obtained the following metrics: precision = 0.784, recall = 0.939, and mAP50 = 0.897.

```
Closing simulation.
(p54) (base) diniththejana@Diniths-MacBook-Air P54-Embodied-Multimodal-LLM-for-Industrial-Task-Planning % python scripts/verify_yolo_workspace.py --
weights yolov8n_p54_workspace.pt --frames 5 --confidence 0.25
pybullet build time: Oct 21 2025 17:04:55
Yolo workspace verification
Repo      : /Users/diniththejana/Documents/GitHub/P54-Embodied-Multimodal-LLM-for-Industrial-Task-Planning
Weights   : yolov8n_p54_workspace.pt
Detector source : yolo
Frames evaluated : 5

Raw YOLODetector.detect(frame) detections:
red block    source=yolo    conf=0.652 pos=(0.45, -0.2, 0.05) bbox={'x_min': 286, 'y_min': 285, 'x_max': 300, 'y_max': 298}
blue block   source=yolo    conf=0.500 pos=(0.25, 0.35, 0.05) bbox={'x_min': 234, 'y_min': 141, 'x_max': 249, 'y_max': 156}
green block  source=yolo    conf=0.343 pos=(0.35, 0.12, 0.05) bbox={'x_min': 259, 'y_min': 280, 'x_max': 275, 'y_max': 214}
yellow block source=yolo    conf=0.284 pos=(0.25, -0.35, 0.05) bbox={'x_min': 234, 'y_min': 322, 'x_max': 250, 'y_max': 338}
left tray    source=yolo    conf=0.978 pos=(0.65, 0.45, 0.01) bbox={'x_min': 320, 'y_min': 100, 'x_max': 373, 'y_max': 143}
right tray   source=yolo    conf=0.735 pos=(0.65, -0.45, 0.01) bbox={'x_min': 321, 'y_min': 336, 'x_max': 372, 'y_max': 379}
workstation  source=yolo    conf=0.921 pos=(0.8, 0.0, 0.0)  bbox={'x_min': 330, 'y_min': 191, 'x_max': 454, 'y_max': 287}

Detected required : ['red block', 'blue block', 'green block', 'yellow block', 'left tray', 'right tray', 'workstation']
Missing required  : []
RESULT: PASS - raw YOLO detected every required workspace object.
(p54) (base) diniththejana@Diniths-MacBook-Air P54-Embodied-Multimodal-LLM-for-Industrial-Task-Planning %
```

Figure : Direct YOLOv8 workspace detection verification. The fine-tuned YOLOv8 model successfully detected all seven required workspace object classes without counting ground-truth fallback as YOLO detection.

SIM-FINAL - PyBullet Execution Path Implementation from Start to Finish:

In addition, I also contributed to the development of the PyBullet execution path implementation process using the KUKA IIWA robot. The KUKA robot controller was connected to the execution interface, and problems associated with execution timeout were solved. The end product was able to receive a natural language pick and place command, form a five-step plan, and execute the LOCATE -> MOVE -> PICK -> MOVE -> PLACE action sequence in the PyBullet environment.




Figure: KUKA IIWA executing the pick-and-place task in PyBullet. The robot executed the generated task plan within the simulated workspace.




Figure: Successful end-to-end pipeline execution. The natural-language instruction was processed through the LLM parser, vision, task planning, execution, and validation stages, with the complete pick-and-place task successfully executed in PyBullet.

INT-TEST - Final Integration and Regression Tests:

Having done the vision and simulation integrations, I did the final integration and regression tests to ensure that the introduced changes did not impact the work of the existing planner, executor and LLM. All the tests from the pytest suite were run and all of them were passed (145

tests passed, 0 fails). The real YOLO detector was tested as well without counting the ground truth fallback as successful YOLO detection. Lastly, the complete real vision and KUKA pipeline was tested successfully.

These contributions supported the Sprint 4 goal of progressing the project from the Alpha prototype stage towards the Beta stage.

```
(p54) (base) diniththejana@Diniths-MacBook-Air P54-Embodied-Multimodal-LLM-for-Industrial-Task-Planning % conda activate p54
pytest -v
cachedir: .pytest_cache
rootdir: /Users/diniththejana/Documents/GitHub/P54-Embodied-Multimodal-LLM-for-Industrial-Task-Planning
configfile: pytest.ini
testpaths: .
plugins: langsmith-0.11.0, anyio-4.14.2
collected 145 items

tests/integration_tests.py::TestSpatialRelationPlanning::test_left_of_applies_negative_x_offset PASSED [ 0%]
tests/integration_tests.py::TestSpatialRelationPlanning::test_right_of_applies_positive_x_offset PASSED [ 1%]
tests/integration_tests.py::TestSpatialRelationPlanning::test_near_applies_diagonal_offset PASSED [ 2%]
tests/integration_tests.py::TestSpatialRelationPlanning::test_spatial_offset_calculation_left PASSED [ 2%]
tests/integration_tests.py::TestSpatialRelationPlanning::test_spatial_offset_calculation_right PASSED [ 3%]
tests/integration_tests.py::TestSpatialRelationPlanning::test_spatial_offset_none_returns_original PASSED [ 4%]
tests/integration_tests.py::TestSpatialRelationPlanning::test_spatial_offset_unknown_relation_uses_default PASSED [ 4%]
tests/integration_tests.py::TestSpatialRelationPlanning::test_in_relation_uses_container_position PASSED [ 5%]
tests/integration_tests.py::TestSpatialRelationPlanning::test_spatial_plan_executes_successfully PASSED [ 6%]
tests/integration_tests.py::TestMultiStepPlanning::test_two_instructions_combined PASSED [ 6%]
tests/integration_tests.py::TestMultiStepPlanning::test_step_numbers_are_sequential PASSED [ 7%]
tests/integration_tests.py::TestMultiStepPlanning::test_multi_step_instruction_preserved PASSED [ 8%]
tests/integration_tests.py::TestMultiStepPlanning::test_multi_step_executes_fully PASSED [ 8%]
tests/integration_tests.py::TestFullPipelineIntegration::test_simple_pick_place_end_to_end PASSED [ 9%]
tests/integration_tests.py::TestFullPipelineIntegration::test_locate_instruction_end_to_end PASSED [ 10%]
tests/integration_tests.py::TestFullPipelineIntegration::test_move_instruction_end_to_end PASSED [ 11%]
tests/integration_tests.py::TestFullPipelineIntegration::test_unknown_object_fails_gracefully PASSED [ 11%]
tests/integration_tests.py::TestFullPipelineIntegration::test_tracker_logs_all_stages PASSED [ 12%]
tests/integration_tests.py::TestFullPipelineIntegration::test_pipeline_robot_state_after_pick_place PASSED [ 13%]
tests/integration_tests.py::TestFullPipelineIntegration::test_consecutive_tasks_independent_state PASSED [ 13%]
tests/integration_tests.py::TestFullPipelineIntegration::test_full_demo_sequence PASSED [ 14%]
tests/integration_tests.py::TestBaselineParser::test_simple_pick_parsed PASSED [ 15%]
tests/integration_tests.py::TestBaselineParser::test_simple_place_parsed PASSED [ 15%]
tests/integration_tests.py::TestBaselineParser::test_synonym_not_handled PASSED [ 16%]
tests/integration_tests.py::TestBaselineParser::test_spatial_not_resolved PASSED [ 17%]
tests/integration_tests.py::TestBaselineParser::test_empty_instruction_fails_gracefully PASSED [ 17%]
tests/integration_tests.py::TestBaselineParser::test_latency_is_sub_millisecond PASSED [ 18%]
tests/integration_tests.py::TestBaselineParser::test_baseline_evaluation_runs PASSED [ 19%]
tests/integration_tests.py::TestBaselineParser::test_baseline_accuracy_lower_than_expected_for_spatial PASSED [ 20%]

PROBLEMS OUTPUT DEBUG CONSOLE TERMINAL PORTS
(p54) diniththejana@Diniths-MacBook-Air P54-Embodied-Multimodal-LLM-for-Industrial-Task-Planning % conda activate p54
pytest -v
tests/test_llm_module.py::TestParserIntegration::test_model_dump_has_all_fields PASSED [ 73%]
tests/test_real_vision_adapter.py::test_get_current_scene_uses_live_simulation_and_returns_json_scene PASSED [ 73%]
tests/test_sprint2.py::TestActionSchema::test_robot_command_valid PASSED [ 74%]
tests/test_sprint2.py::TestActionSchema::test_position_as_tuple PASSED [ 75%]
tests/test_sprint2.py::TestActionSchema::test_action_plan_total_steps PASSED [ 75%]
tests/test_sprint2.py::TestActionSchema::test_command_summary PASSED [ 76%]
tests/test_sprint2.py::TestActionSchema::test_plan_to_commands_simple PASSED [ 77%]
tests/test_sprint2.py::TestActionSchema::test_plan_to_commands_with_task_id PASSED [ 77%]
tests/test_sprint2.py::TestActionSchema::test_all_command_types_valid PASSED [ 78%]
tests/test_sprint2.py::TestMockRobot::test_scene_loads_correctly PASSED [ 79%]
tests/test_sprint2.py::TestMockRobot::test_move_to_valid_position PASSED [ 80%]
tests/test_sprint2.py::TestMockRobot::test_move_to_out_of_bounds PASSED [ 80%]
tests/test_sprint2.py::TestMockRobot::test_move_to_object PASSED [ 81%]
tests/test_sprint2.py::TestMockRobot::test_move_to_unknown_object PASSED [ 82%]
tests/test_sprint2.py::TestMockRobot::test_pick_success PASSED [ 82%]
tests/test_sprint2.py::TestMockRobot::test_pick_unknown_object PASSED [ 83%]
tests/test_sprint2.py::TestMockRobot::test_pick_when_already_holding PASSED [ 84%]
tests/test_sprint2.py::TestMockRobot::test_place_success PASSED [ 84%]
tests/test_sprint2.py::TestMockRobot::test_place_without_holding PASSED [ 85%]
tests/test_sprint2.py::TestMockRobot::test_place_unknown_location PASSED [ 86%]
tests/test_sprint2.py::TestMockRobot::test_locate_success PASSED [ 86%]
tests/test_sprint2.py::TestMockRobot::test_locate_unknown_object PASSED [ 87%]
tests/test_sprint2.py::TestMockRobot::test_full_pick_place_sequence PASSED [ 88%]
tests/test_sprint2.py::TestMockRobot::test_robot_reset PASSED [ 88%]
tests/test_sprint2.py::TestMockRobot::test_command_log_grows PASSED [ 89%]
tests/test_sprint2.py::TestMockRobot::test_case_insensitive_object_lookup PASSED [ 90%]
tests/test_sprint2.py::TestExecutor::test_simple_locate_plan_succeeds PASSED [ 91%]
tests/test_sprint2.py::TestExecutor::test_full_pick_place_plan_succeeds PASSED [ 91%]
tests/test_sprint2.py::TestExecutor::test_plan_fails_on_unknown_object PASSED [ 92%]
tests/test_sprint2.py::TestExecutor::test_plan_stops_on_first_failure PASSED [ 93%]
tests/test_sprint2.py::TestExecutor::test_execution_result_has_latency PASSED [ 93%]
tests/test_sprint2.py::TestTaskPlanner::test_pick_plan_generates_steps PASSED [ 94%]
tests/test_sprint2.py::TestTaskPlanner::test_place_plan_generates_full_sequence PASSED [ 95%]
tests/test_sprint2.py::TestTaskPlanner::test_locate_plan_is_single_step PASSED [ 95%]
tests/test_sprint2.py::TestTaskPlanner::test_pick_with_destination_includes_place PASSED [ 96%]
tests/test_sprint2.py::TestTaskPlanner::test_missing_object_raises PASSED [ 97%]
tests/test_sprint2.py::TestTaskPlanner::test_plan_preserves_instruction PASSED [ 97%]
tests/test_sprint2.py::TestTaskPlanner::test_plan_with_task_id PASSED [ 98%]
tests/test_sprint2.py::TestTaskPlanner::test_move_plan_same_as_place PASSED [ 99%]
tests/test_sprint2.py::TestTaskPlanner::test_full_pipeline_planner_to_executor PASSED [100%]

===== 145 passed in 30.48s =====
(p54) diniththejana@Diniths-MacBook-Air P54-Embodied-Multimodal-LLM-for-Industrial-Task-Planning %
```

Figure: Final integration and regression test results. The complete pytest suite successfully passed 145 tests with no failures, confirming that the final integration changes did not introduce regressions into the existing system.

Kaveesha Dulanjana Dharmadasa (105271678)

My contribution this sprint was documentation. I took EPIC 3, the final report, and REP-4, the discussion and conclusion.

REP-4, discussion. I wrote the complete discussion section of the final report. Rather than working from the summary tables in the Sprint 3 evaluation report, I read `llm_backend/LLM_eval/metrics.py`, `baseline_parser.py` and `task_planner/planner.py` first, so that every claim in the discussion is tied to what the committed code and data actually show.

The section argues the central research claim, that the accuracy gap between GPT-4o and the rule based parser measures the value of putting an LLM in the loop. Across the 25 case run GPT-4o reaches 76.0 percent instruction accuracy against the baseline's 52.0 percent, with parse success at 100.0 against 88.0 and an error rate of 0.0 against 12.0.

The more useful result is that the gap is not uniform. GPT-4o wins decisively on simple instructions, 100.0 against 60.0, and on spatial instructions, 100.0 against 60.0, and roughly doubles the baseline on multi step, 66.7 against 33.3. It adds nothing on synonym instructions, 60.0 against 60.0, and nothing on edge cases, 75.0 against 75.0. Both null results are written out rather than left in a table. On ambiguous instructions GPT-4o scores 33.3 percent, which is poor in absolute terms, but it records 100.0 percent action accuracy and 100.0 percent confidence calibration on the same cases, meaning it identifies the intended action and correctly reports low confidence rather than guessing. For a system that drives a robot arm that distinction matters more than the headline number.

The limitations subsection covers the small per category sample, the single location latency measurement, the absence of any hardware trial, the detector's bias toward false positives given precision of 0.784 against recall of 0.939, and the loss of pre-execution boundary validation described below.




Figure 8. Instruction accuracy by category, GPT-4o against the rule based baseline, from evaluation_metrics.csv and eval_report.csv.




Figure 9. Summary statistics for GPT-4o parse latency across the ten instruction benchmark, from the output of benchmark.py.

REP-4, conclusion. I wrote the conclusion to answer the research question rather than restate the abstract, and to give an honest read on deployment feasibility. Since all validation is in simulation, the claim the report makes is that the architecture is ready for a hardware trial, not that it is ready for a factory floor.

EPIC 3, report assembly and consistency pass. This was self-assigned. I did a full pass over the assembled document and checked every quantitative claim against its source in the repository. Three findings changed what the report claims.

First, the figures in P54_Evaluation_Report.pdf do not reconcile with the committed data. That report quotes GPT-4o at 85.0 percent overall instruction accuracy, 88.0 percent on ambiguous and 83.3 percent on multi step, where evaluation_metrics.csv gives 76.0, 33.3 and 66.7. The PDF figures are also unreachable given the dataset, since the ambiguous and multi step categories hold three cases each and 88.0 percent is not a possible value over three cases. The final report uses the committed CSV throughout.

Second, the nine primitive decomposition described in the Sprint 2 documentation is no longer what the planner produces. task_planner/planner.py emits a five command plan of LOCATE, MOVE, PICK, MOVE, PLACE for a pick and place instruction with a destination, and shorter plans for a bare location. Five commands is what Dinith's execution section describes. The methodology had inherited the older description and has been corrected.

Third, the Sprint 2 safety interceptor no longer exists. task_planner/safety.py, which raised a BoundaryViolation for coordinates outside X in [0.1, 0.8], Y in [-0.4, 0.4] or Z in [0.0, 0.5], was lost during the planner rewrite. The remaining checks are Kuka_IIWA._within_bounds() with wider limits and a 2D workspace check in MockRobot, both of which run inside the execution layer rather than before it. The report no longer claims a pre-execution guardrail, and the gap has been added to the risk register.

Routine reconciliation covered the dataset size, quoted as 25 cases in the Sprint 3 report and 35 in the current methodology; the regression count, reconciled between 142 verifications in Sprint 3 and the current 145 test functions; terminology drift between "primitive", "robot command" and "action step"; and figure and heading numbering across the merged sections.

Minh Hoang Duong

During Sprint 4, my contribution centred on repository hygiene and a structural refactor of the codebase, followed by

CODE-1 — Clean the repository and remove unnecessary/dead files

I removed dead code left over from earlier development — the unused `drafts/vision_backend/` stub package, a duplicate `__init__.py`, an empty `Universal_Robot_ur5.py` stub — and archived the old report-generation draft rather than deleting it outright. I also added `reportlab` to `requirements.txt`, a dependency that had been missing. This reduced clutter and stale code paths that could otherwise confuse teammates picking up the vision or reporting modules.

Codebase refactor

I restructured the repository so its layout matches how the pipeline actually runs — consolidating the vision entry point (`scene_representation.py`) into `simulation_backend/vision/`, relocating detector weights and evaluation artefacts, and updating `README.md` and docstrings to match.

Fine-tuning reorganisation and integrity audit.

I moved the YOLO fine-tuning dataset, weights, and helper scripts into dedicated `fine_tuning/` and `helper_scripts/` directories, and traced every path reference the moved scripts relied on. This caught several breakages the move alone would have introduced silently: incorrect repo-root resolution in both fine-tuning scripts, stale default paths, and a teammate's machine-specific absolute path baked into `args.yaml/data.yaml`.

I also found and fixed the PDF report generator, which was accepting `--csv/--json` arguments but never actually reading them — it always rendered fixed placeholder figures regardless of the real evaluation run. This directly protects the accuracy of the final report and the reproducibility of Dinith's fine-tuning work on other machines.

Vision detector wiring problem

Auditing weight references, I found a duplicate `yolov8n.pt` in both the repo root and `simulation_backend/vision/detection_weight`, and removed the root copy. This exposed a deeper issue: with `VISION_DETECTOR=yolo`, the pipeline loads the stock COCO-pretrained checkpoint rather than Dinith's fine-tuned model, so it can't recognise our custom workspace labels and silently falls back to ground truth. This is not yet fixed




Figure 10. The yolo8 detection problems in pipeline, falling back to ground_truth.

4. SPRINT REVIEW (CRITICAL REVIEW OF THE PRODUCT)

Demonstration Session

The Sprint 4 demonstration presented the project as a single integrated Beta-stage system rather than as separate technical modules. The team gave a brief overview of the system architecture, then ran representative natural-language instructions end to end: LLM parsing, live scene perception, task planning, and PyBullet execution on the KUKA IIWA arm.

Three categories of scenario were demonstrated. First, simple and spatial pick-and-place instructions (e.g. "pick up the red block and place it on the blue tray", "place the block behind the workstation") were shown succeeding end to end, matching the 100% simple-instruction and 100% spatial-instruction accuracy recorded in the Sprint 4 evaluation (Section "Empirical Performance Summary" below). Second, a multi-step instruction requiring sequential pick-and-place actions was demonstrated to show the planner's `plan_multi_step()` decomposition. Third, two failure-handling scenarios were run deliberately: an instruction referencing a non-existent object, and an ambiguous instruction with no object specified, to show that the system rejects or flags low-confidence instructions rather than executing an unsafe or arbitrary action.

The completed pytest regression suite (145/145 tests passing) and the fine-tuned YOLOv8 detection results (precision 0.784, recall 0.939, mAP50 0.897) were presented as supporting evidence that the demonstrated behaviour reflects the state of the underlying codebase rather than a one-off manually staged run.

Client Feedback

The client's response to the Sprint 4 demonstration was strongly positive. The client was pleased with how far the system has progressed since the Alpha prototype, and gave positive feedback on the reliability shown during the live end-to-end run, including the successful pick-and-place scenarios and the system's handling of the deliberately staged failure cases.

Overall client sentiment: Satisfied. The client expressed satisfaction with the progress made this sprint and the overall direction of the project, and did not raise any concerns about the demonstrated functionality.

Empirical Performance Summary

Table 1 consolidates the Sprint 4 evaluation results (evaluation_metrics.csv, 35-case dataset, GPT-4o vs the rule-based baseline parser) alongside the operational metrics gathered during integration testing and benchmarking.

Metric	GPT-4o (LLM)	Rule-based baseline	Note
Overall instruction accuracy	76.0%	52.0%	24-point gap; driven almost entirely by spatial and multi-step categories
Parse success rate	100.0%	88.0%	Baseline fails to produce any structured output for 12% of instructions
Pipeline error rate	0.0%	12.0%	See “Why the error rate is 0%” below — this is a structural-validity metric, not a correctness metric
Simple instructions	100%	60%	Baseline handles literal keyword matches only
Spatial instructions	100%	60%	Baseline has no representation of relative position (“behind”, “left of”)
Multi-step instructions	66.7%	33.3%	Baseline cannot decompose compound instructions into ordered sub-actions
Synonym instructions	60%	60%	No LLM advantage — both approaches handle known synonyms equally
Edge-case instructions	75%	75%	No LLM advantage on this category in the current 35-case set
Ambiguous instructions	33.3%	0%	GPT-4o records 100% action accuracy and 100% confidence calibration on these cases — see discussion below
Mean / P95 latency	1,433 ms / 2,461 ms	n/a (rule-based, sub-millisecond)	Latency is the cost of the accuracy gain, not a defect
Vision detector (fine-tuned YOLOv8)	Precision 0.784 / Recall 0.939 / mAP50 0.897	n/a	Detector biased toward false positives (lower precision than recall)
Regression test suite	145 / 145 passed	n/a	Confirms Sprint 4 changes did not break existing planner/executor/LLM behaviour

Table 1. Sprint 4 empirical performance summary: GPT-4o vs rule-based baseline, from evaluation_metrics.csv, benchmark.py output, and the Sprint 4 pytest run.

Critical Analysis of Progress and Results

Sprint 4 progressed the project from the Alpha prototype towards a more stable Beta-stage system, and the evaluation results in Table 1 give a way to judge that claim quantitatively rather than descriptively. The headline 76.0% vs 52.0% instruction-accuracy gap is not evenly distributed, and the distribution matters more than the average: GPT-4o wins decisively wherever the instruction requires semantic or spatial reasoning the rule-based parser cannot represent (spatial, multi-step), and shows no advantage at all where both approaches reduce to lookup (synonym, edge case). This is the expected signature of a genuine capability gap rather than an artefact of prompt tuning — if the LLM were simply memorising surface patterns, its advantage would be roughly uniform across categories instead of concentrated where the baseline's fixed grammar structurally cannot cope.

The rule-based baseline's failures are explainable rather than incidental. It matches instructions against a fixed set of keyword and regular-expression patterns for object, action and

destination. It has no mechanism for representing relative spatial relations (“behind”, “to the left of”), so spatial instructions either fail to parse or default to an incorrect literal interpretation — hence 60% rather than a near-zero score, since some spatial instructions happen to also contain a literal keyword match. It has no compositional structure for chaining sub-actions, so multi-step instructions are only handled when they happen to reduce to a single recognisable pattern, giving the 33.3% multi-step result. This is a scientific explanation of the failure mode, not just an observation of a lower number: the baseline is architecturally incapable of the categories where it underperforms, which is precisely the comparison the sprint's research question (“what does putting an LLM in the loop actually buy you?”) was designed to isolate.

The ambiguous-instruction result deserves separate treatment because the headline number (33.3%) looks like a weakness but the supporting metrics show the opposite. On these cases GPT-4o still achieves 100% action accuracy and 100% confidence calibration, meaning it correctly identifies what the user is likely asking for and correctly reports that its confidence is too low to proceed, rather than guessing and executing an incorrect action. For a system whose output drives a physical robot arm, correctly flagging “I don't have enough information to act safely” is a more valuable behaviour than a higher raw accuracy score achieved by guessing, and the current metric (instruction accuracy) does not reward this distinction. This is flagged as a methodological limitation for the final report: a single instruction-accuracy number understates the system's real safety-relevant performance on ambiguous input.

How the System is Safe

Safety in this pipeline is treated as a property that has to be actively engineered at multiple stages, not an emergent property of the LLM being “good”. Three concrete mechanisms currently provide protection, and one previously-planned mechanism is currently missing and has been logged as a risk rather than left unstated.

- Schema-level validation: every LLM output is constrained to a Pydantic v2 schema before it is accepted. Malformed or incomplete responses are rejected at the parser boundary and never reach the planner, which is the direct reason the pipeline records a 0% error rate (see below).
- Confidence-gated execution: when the parser's confidence score for an interpretation is low (as with ambiguous instructions), the system is designed to flag the instruction rather than execute it, which is why the ambiguous category shows 100% confidence calibration alongside a low raw accuracy score.
- Execution-layer boundary checks: `Kuka_IIWA._within_bounds()` and the `MockRobot 2D` workspace check reject commanded positions outside the physically safe workspace before a movement is issued to the simulated (and, in a future hardware trial, real) robot arm.

- Known gap: the Sprint 2 pre-execution safety interceptor (`task_planner/safety.py`, which raised a `BoundaryViolation` for coordinates outside a defined X/Y/Z envelope) was lost during the planner rewrite and has not been reinstated. The remaining checks above run inside the execution layer rather than strictly before it, which is a narrower safety margin than the architecture originally specified. This gap has been added to the Sprint 4 risk register rather than omitted from the report.

Taken together, the system is safe in the specific sense that invalid or unrecognisable instructions cannot reach the robot as commands (schema validation), that low-confidence interpretations are distinguishable from confident ones (confidence gating), and that commanded coordinates are checked against a workspace envelope before movement (boundary checks) — but it is not yet safe in the stronger sense of having an independent pre-execution interceptor that is decoupled from the execution layer itself, which is the outstanding item for the Beta hardening work.

Why the Reported Error Rate is 0%

The 0.0% figure in Table 1 is easy to misread as “the system is never wrong”, which is not what it measures. “Error rate” in the evaluation harness counts pipeline-level failures — exceptions, unparseable output, or a response that does not satisfy the output schema — not incorrect answers. GPT-4o's 0.0% error rate means that, across all 25/35 evaluated instructions, the model always returned a structurally valid, schema-compliant response; it does not mean every response was correct, and Table 1's 76.0% instruction-accuracy figure is the metric that captures correctness. The two numbers together describe two different properties: GPT-4o is 100% reliable at producing a well-formed answer and 76.0% accurate at producing the right answer, and it is the gap between those two figures — not either number alone — that describes its real behaviour.

The baseline's 12.0% error rate is the more informative comparison point. The rule-based parser has no fallback when an instruction does not match any of its patterns: it either matches and returns a parse, or it fails outright and raises an exception that the harness counts as an error. GPT-4o structurally cannot fail this way because it always produces some schema-valid completion (assisted by prompt design and the Pydantic validation layer), even when that completion is semantically wrong. This is a genuine engineering property of the schema-first design (Section 3.2 of the final report methodology) rather than a claim about model correctness, and the final report should state both numbers together rather than quoting the 0% figure in isolation.

Critical Analysis of the Process

The main technical challenge of Sprint 4 was reliability at module boundaries rather than within any single module: the planner, vision adapter and executor each passed their own unit and component tests, but integration failures still occurred where the live scene representation's

object/position schema drifted from what the planner expected. This is consistent with the EPIC 3 consistency-pass findings — the nine-primitive decomposition documented in Sprint 2 no longer matches what the planner actually emits (a five-command LOCATE/MOVE/PICK/MOVE/PLACE plan), and the Sprint 2 safety interceptor referenced in earlier documentation no longer exists in the codebase. Both discrepancies were caught by re-deriving report claims directly from the committed source and data rather than carrying forward descriptions from earlier sprint reports, which is the process change the team is recommending going into the remaining sprints.

Overall, the sprint moved the project closer to the Beta objective by prioritising validation, failure handling and evidence traceability over new feature surface area. The remaining work for the next sprint is to close the pre-execution safety-interceptor gap and expand the ambiguous/multi-step evaluation categories, which currently rest on only three cases each and are too small to support a strong statistical claim.

5. RETROSPECT (CRITICAL REVIEW OF THE PROCESS)

The team continued to follow an Agile development process using Jira for task management, GitHub for version control, and regular team and supervisor meetings. The Sprint 4 work was divided according to the existing module ownership while integration issues were discussed collectively.

A major strength of the sprint was that the team entered Semester 2 with a functioning Alpha system rather than having to rebuild the core architecture. This allowed the team to focus on improving reliability and testing instead of implementing the basic components again.

One challenge was that changes to one module could affect other parts of the pipeline. For example, changes to the live scene representation could affect the planner because the planner depends on a specific object and position format. This reinforced the importance of maintaining shared interfaces and testing module boundaries.

Another challenge was the development environment. The project uses PyBullet and several machine-learning dependencies, which can behave differently across operating systems and Python environments. The team addressed this through documented dependencies, environment configuration and testing on supported systems.

From a cybersecurity and ethical perspective, the team continues to avoid using confidential industrial data during development. The LLM component can communicate with external providers such as OpenAI, Gemini and DeepSeek, meaning that real industrial instructions could potentially expose sensitive information if used without appropriate controls. The project therefore continues to treat local LLM inference as an important option for future deployment.

The team also considered the safety implications of using an LLM to generate robot actions. Invalid or incorrect instructions should not directly result in robot movement. Existing structured schemas, object validation and workspace safety checks provide layers of protection, while Sprint 4 testing focuses on identifying additional failure cases.

Overall, the team process was effective in supporting incremental development, although the sprint highlighted the need for stronger regression testing whenever changes are made to shared interfaces.

6. LESSONS LEARNED (CRITICAL REVIEW OF SPRINT 4 EXPERIENCE AND FUTURE PLAN)

- **Integration needs continuous regression testing.** The existing end-to-end pipeline can be affected by changes to individual modules, so integration testing should be performed throughout development rather than only before the final demonstration.
- **Shared interfaces must remain stable.** The live vision system and task planner depend on a consistent scene structure. Changes to object labels or coordinate formats can create failures even when both modules work independently.
- **Successful execution is not enough for a Beta system.** The system must also respond safely when objects are missing, instructions are invalid or execution fails. Future testing should therefore include both successful and failure scenarios.
- **Environment configuration affects reproducibility.** PyBullet and machine-learning dependencies can create problems across different machines. A consistent setup and clear installation instructions should be maintained throughout the remaining development.
- **Testing should reflect realistic user instructions.** Simple pick-and-place commands are useful for demonstrating functionality, but the system should also be tested with spatial, multi-step, ambiguous and invalid instructions.
- **Safety should be considered at every pipeline stage.** Since the final output controls simulated robot behaviour, validation should occur before commands reach the execution layer rather than relying only on the robot simulator to prevent unsafe actions.

Recommendations for the next sprint

1. Expand end-to-end testing using a larger set of realistic task instructions.
2. Resolve remaining runtime and environment issues before the next client demonstration.
3. Continue improving execution failure handling and recovery.
4. Measure end-to-end latency and task success rate.
5. Finalise the Beta test results and use them to identify the remaining issues before final development and handover.

The lessons are aligned with the Semester 1 retrospective, which identified integration testing, shared interfaces, environment compatibility and systematic evaluation as key areas for improvement.

7. APPENDIX A: PSEUDOCODE OF KEY PIPELINE COMPONENTS

The following pseudocode summarises the structure of the five core pipeline components described in the Sprint Progress and methodology sections. It is a simplified, language-agnostic representation of the committed implementation intended to make the control flow and safety checks auditable without reproducing the full source — the authoritative implementation is in the referenced files in the project repository.

A.1 Pipeline Orchestrator (main.py)

```
FUNCTION run_pipeline(instruction_text, scene_source):
    scene = vision.get_live_scene(scene_source) # A.4
    parsed_instruction = llm_parser.parse(instruction_text, scene) # A.2
    IF parsed_instruction.confidence < CONFIDENCE_THRESHOLD:
        RETURN Result(status=REJECTED, reason='low_confidence')
    plan = task_planner.plan(parsed_instruction, scene) # A.3
    IF plan IS INVALID:
        RETURN Result(status=REJECTED, reason='no_valid_plan')
    result = executor.execute(plan) # A.5
    RETURN result
```

A.2 LLM Instruction Parser (schema-first)

```
FUNCTION parse(instruction_text, scene):
    prompt = build_prompt(instruction_text, scene, few_shot_examples)
    raw_response = call_llm_backend(prompt) # GPT-4o / Gemini / DeepSeek /
    local
    TRY:
        parsed = ParsedInstructionSchema.validate(raw_response) # Pydantic
    v2
    CATCH SchemaValidationError:
        RETURN ParsedInstruction(confidence=0, status=INVALID) # never
    raises
    RETURN parsed # always a structurally valid object -> 0% pipeline error
    rate
```

A.3 Task Planner (rule-based decomposition)

```
FUNCTION plan(parsed_instruction, scene):
    IF parsed_instruction.is_multi_step:
        RETURN plan_multi_step(parsed_instruction, scene)
    target_object = resolve_object(parsed_instruction.object, scene)
    IF target_object IS NULL:
        RETURN INVALID_PLAN
    destination = resolve_destination(parsed_instruction.destination,
    scene,
    SPATIAL_OFFSETS)
    RETURN [LOCATE(target_object), MOVE(target_object),
    PICK(target_object), MOVE(destination), PLACE(destination)]
```

A.4 Vision / Scene Representation (with graceful fallback)

```
FUNCTION get_live_scene(scene_source):
    IF VISION_DETECTOR == 'yolo':
        detections = yolo_model.detect(scene_source.camera_frame)
        IF detections.confidence_ok AND
        labels_match_workspace_classes(detections):
            RETURN build_scene(detections)
        LOG_WARNING('vision fallback triggered')
    RETURN build_scene(scene_source.ground_truth) # fallback, flagged in
Fig.10
```

A.5 Execution Layer (boundary-checked)

```
FUNCTION execute(plan):
    FOR step IN plan:
        IF step.type == MOVE:
            IF NOT robot._within_bounds(step.target_position): # safety
                check
                RETURN Result(status=ABORTED, reason='out_of_bounds')
            TRY:
                robot.perform(step)
            CATCH ExecutionTimeout:
                RETURN Result(status=FAILED, reason='timeout', step=step)
    RETURN Result(status=SUCCESS, trace=execution_log)
```

Note: task_planner/safety.py (the standalone pre-execution BoundaryViolation interceptor from Sprint 2) is not shown here because it no longer exists in the current codebase — see “How the System is Safe” for the resulting risk.

APPENDIX B: REFERENCES

- Ahn, M., Brohan, A., Brown, N., Chebotar, Y., Cortes, O., David, B., Finn, C., Fu, C., Gopalakrishnan, K., Hausman, K., et al. (2022). Do As I Can, Not As I Say: Grounding Language in Robotic Affordances. arXiv preprint arXiv:2204.01691.
- Driess, D., Xia, F., Sajjadi, M. S. M., Lynch, C., Chowdhery, A., Ichter, B., Wahid, A., Tompson, J., Vuong, Q., Yu, T., et al. (2023). PaLM-E: An Embodied Multimodal Language Model. arXiv preprint arXiv:2303.03378.
- Liang, P., Bommasani, R., Lee, T., Tsipras, D., Soylu, D., Yasunaga, M., Zhang, Y., Narayanan, D., Wu, Y., Kumar, A., et al. (2022). Holistic Evaluation of Language Models. arXiv preprint arXiv:2211.09110.